



University
of Glasgow

A WORLD
TOP 100
UNIVERSITY

Introduction to Data Science and Systems

Lecture 1: Vectors and Matrices

Dr Iraklis Klampanos

Dr Nicolas Pugeault

WORLD
CHANGING
GLASGOW

THE TIMES
THE SUNDAY TIMES
GOOD
UNIVERSITY
GUIDE
2024
SCOTTISH
UNIVERSITY
OF THE YEAR




Intended Learning Outcomes

- what are vectors and vector spaces
- the standard operations on vectors: addition and multiplication
- what a norm is and how it can be used to measure vectors
- what an inner product is and how it gives rise to geometry of vectors
- how mathematical vectors map onto numerical arrays
- the different p-norms and their uses
- important computational uses of vector representations
- how to characterise vector data with a mean vector and a covariance matrix
- the properties of high-dimensional vector spaces
- the basic notation for matrices
- the view of matrices as linear maps
- how basic geometric transforms are implemented using matrices
- how matrix multiplication is defined and its algebraic properties
- the basic anatomy of matrices



Example: Text and translation

- Text, as represented by strings in memory, has *weak structure*.
- There are **comparison functions** for strings (e.g. edit distance, Hamming distance), but **only character-level semantics**
- **String operations** are character-level operations like concatenation or reversal, but not useful for machine translation system






Words aren't enough

- *Original*
 - "The craft held fast to **the bank of the burn.**"
 - (*the vessel stayed moored to **the edge of the stream***)
- *Dictionary lookup (naïve)*
 - French: "L'artisanat tenu rapide à la Banque de la brûlure."
 - (*the artisanal skill held quickly to **the financial institution of the burn wounds***)
 - Danish: "Håndværket holdt fast til banken af brænden"
 - (*The craftsmanship held fast **to the bank [financial institution] of fire***)
- *Correct(ish)*
 - French: "Le bateau se tenait fermement à la rive du ruisseau."
 - (*the boat was firmly attached to the **riverbank***)
 - Danish: "Farttøjet holdt fast i bredden af åen"
 - (*The vessel held fast at **the bank of the river***)




Solution -- to place them in a vector space


- Imbue text fragments with additional mathematical structure -- **to place them in a vector space**
- Fragments might be words, partial words or whole sentences



Male-Female



Verb tense



Country-Capital



The structure of a (topological) vector space

- What words are like **salamander**?
 - **Distance/metric functions:** **norm**
 - E.g. the neighbourhood of the vector corresponding to salamander, which might include words like axolotl or waterdog
- What is the equivalent of a king, but with a woman instead of a man?
 - **Operation functions:** **subtraction, mean**
 - Famously, the original word2vec paper notes that on test data, the equation

$$\text{King} - \text{Man} + \text{Woman} = \text{Queen}$$



Vector spaces

- **Vectors** to be ordered tuples of real numbers

$$[x_1, x_2, \dots, x_n], x_i \in \mathbb{R}$$

- A vector has a fixed dimension n
- Each element of the vector as representing a distance in a **direction orthogonal** to all the other elements.
- Orthogonal just means "independent", or, geometrically speaking "at 90 degrees".
- We write vectors with a bold lower case letter:

$$\mathbf{x} = [x_1, x_2, \dots, x_d],$$

$$\mathbf{y} = [y_1, y_2, \dots, y_d],$$



Vector spaces

- For example, a length-3 vector might be used to represent a spatial position in **Cartesian** coordinates, with three orthogonal measurements for each vector.
- Consider the 3D vector $[5, 7, 3]$

$$\begin{aligned} &5 * [1, 0, 0] + \\ &7 * [0, 1, 0] + \\ &3 * [0, 0, 1] \end{aligned}$$

- Each of these vectors $[1, 0, 0]$, $[0, 1, 0]$, $[0, 0, 1]$ is pointing in an independent direction (orthogonal direction) and has length one.



Points in space

Notation:

- $\mathbb{R}$ denotes the **set of real numbers**.
- $\mathbb{R}_{\geq 0}$ denotes the set of non-negative **reals**.
- $\mathbb{R}^n$ denotes the set of tuples of exactly n real numbers (**vector**).
- $\mathbb{R}^{n \times m}$ denotes the set of **2D arrays (matrix)** of real numbers arranged in n rows by m columns.
- The notation $(\mathbb{R}^n, \mathbb{R}^n) \rightarrow \mathbb{R}$ says that the **operation** defines a map from a pair of n dimensional vectors to a real number.



Points in space

Vector spaces

Any vector of given dimension n lies in a **vector space**, called $\mathbb{R}^n$, along with the operations of:

- **scalar multiplication** so that $a\mathbf{x}$ is defined for any scalar a . For real vectors, $a\mathbf{x} = [ax_1, ax_2, \dots, ax_n]$, elementwise scaling.
 $(\mathbb{R}, \mathbb{R}^n) \rightarrow \mathbb{R}^n$
- **vector addition** so that $\mathbf{x} + \mathbf{y}$ vectors $\mathbf{x}, \mathbf{y}$ of equal dimension. For real vectors, $\mathbf{x} + \mathbf{y} = [x_1 + y_1, x_2 + y_2, \dots, x_d + y_d]$ the elementwise sum
 $(\mathbb{R}^n, \mathbb{R}^n) \rightarrow \mathbb{R}^n$



Points in space

Two additional operations

- **norm** $\| \mathbf{x} \|$ which allows the length of vectors to be measured.

$$\mathbb{R}_n \rightarrow \mathbb{R}_{\geq 0}$$

- **inner product** $\langle \mathbf{x} | \mathbf{y} \rangle$, $\langle \mathbf{x}, \mathbf{y} \rangle$ or $\mathbf{x} \cdot \mathbf{y}$ which determines the **angle** between two vectors. The inner product of two orthogonal vectors is 0 . For real vectors

$$\mathbf{x} \cdot \mathbf{y} = x_1 y_1 + x_2 y_2 + x_3 y_3 \dots x_d y_d$$

$$(\mathbb{R}^n, \mathbb{R}^n) \rightarrow \mathbb{R}$$

Topological and inner product spaces

- With a norm a vector space is a **normed vector space/topological vector space**.
- With an inner product, a vector space is an **inner product space**, and we can talk about the angle between two vectors.

Topological/geometrical space:


- a set whose elements are called points

Metric space:

- define distance between points, e.g. norm


Normed vector space:

- a vector space with norm defined



Topological and inner product spaces

- With a norm a vector space is a **normed vector space**.
- With an inner product, a vector space is an **inner product space**
- Relationship between **Inner Product** and **Norm**:
 - If you have an inner product on a vector space, you can derive a norm from it: $\|\mathbf{x}\| = \sqrt{\langle \mathbf{x} | \mathbf{x} \rangle}$
 - However, not all norms come from an inner product.






Vectors

- Points in space
- Arrows pointing from the origin
- Tuples of numbers

These are all valid ways of thinking about vectors.

The "**points in space**" mental model is probably the most useful

- vectors to represent *data*; data lies in space
- matrices to represent *operations* on data; matrices warp space.



Relation to arrays

- 1D floating point arrays are we called "vectors"
- floating point numbers are **NOT** real numbers

```
# two 3D vectors (3 element ordered tuples)  
x = np.array([0,1,1])  
y = np.array([4,5,6])  
a = 2  
print_matrix("a", a)  
print_matrix("x", x)  
print_matrix("y", y)
```

$a = 2$

x

$\begin{bmatrix} 0 & 1 & 1 \end{bmatrix}$

y

$\begin{bmatrix} 4 & 5 & 6 \end{bmatrix}$



Uses of vectors

They are a *lingua franca* for data. Because vectors can be

- **composed** (via addition),
- **compared** (via norms/inner products)
- and **weighted** (by scaling),

In Numpy, they map onto the efficient **ndarray** data structure, so we can operate on them efficiently and concisely.



Vector data

- Datasets are commonly stored as 2D **tables**.

heart_rate	systolic	diastolic	vo2
67	110	72	98
65	111	70	98
64	110	69	97
..			

Observations

one element of the vector (feature)


Each **row** can be seen as a vector in $\mathbb{R}^n$



Geometric operations

- Use vectors in 3D space
- Transformations in 3D space include:
 - scaling
 - rotation
 - flipping (mirroring)
 - translation (shifting)

```
[[ -0.   15.   0. ]  
[ 16.  -0.5  32.5]  
[-16.  -0.5  32.5]  
[ 16. -15. -32.5]  
[-16. -15. -32.5]  
[-44.  10. -32.5]  
[-60.  -3. -13. ]  
[-65.  -3. -32.5]  
[ 44.  10. -32.5]  
[ 60.  -3. -13. ]  
[ 65.  -3. -32.5]  
[ -0.   15. -32.5]]
```



The *Cobra Mk. III* spaceship model above is defined by these vectors specifying the vertices in 3D space



Machine learning applications


- Machine learning relies heavily on vector representation. A typical machine learning process involves:
 - transforming some data onto **feature vectors**
 - creating a function that transforms **feature vectors** to a prediction (e.g. a class label)
- *Most machine learning algorithms can be seen as doing geometric operations: **comparing distances, warping space, computing angles, and so on.***
 - E.g. **k nearest neighbours**



Example: irises classification

The irises classification task:

- *The measurements of the dimensions of the sepals and petals of irises allows classification of species*
 - **Training data:** [sepal height, petal]
 - **Labels:** species of irises
-
- ***k* nearest neighbours**
 - Using a **norm** to compute distances
 - The output prediction is the class label that occurs most times among these *k* neighbours








Example: irises classification

- **k nearest neighbours**


- Calculate distance between **training examples** and the target using **norm**
- Finding the closest k neighbors (e.g. 3, 5, 10), voting for the **most majority**






Example: irises classification

- *The choice of k might significantly affect the prediction result*
- *The distance function (norm or cosine similarity) is something need to be considered*



Example: Image compression

- Images can be represented as 2D arrays of brightness
- Groups of pixels -- for example, rectangular patches -- can be unraveled into a vector.
 - E.g. An 8x8 image patch would be unraveled to a 64-dimensional vector.
- Splitting images into patches, and treating each patch as a vector $x_1, \dots, x_n$
- The vectors are **clustered** to find a small number of vectors $y_1, \dots, y_m$, $m \ll n$ that are a reasonable approximation of nearby vectors.
- the vectors for the small number of representative vectors y_i are stored (the **codebook**)






Basic vector operations

- Standard operations:
 - getting the length of vectors (**norm**)
 - computing dot (inner), outer and cross products.
- **Addition and multiplication** -- form **weighted sums** of vectors

$$\lambda_1 \mathbf{x}_1 + \lambda_2 \mathbf{x}_2 + \cdots + \lambda_n \mathbf{x}_n$$




See the lecture note for codes



Basic vector operations

- Many standard statistics and operations can be directly applied.
- **Linear interpolation:** to construct **new** data points within the range of **known** data points.

$$\text{lerp}(\mathbf{x}, \mathbf{y}, \alpha) = (1 - \alpha)\mathbf{x} + (\alpha)\mathbf{y}$$



See the lecture note for codes



How big is that vector?

- The L_2 -norm or Euclidean length of a vector x (written as $\|x\|$) can be computed directly using `np.linalg.norm()`.
- This is equal to:

$$\|\mathbf{x}\|_2 = \sqrt{x_0^2 + x_1^2 + x_2^2 + \cdots + x_n^2}$$

```
x = np.array([1.0, 10.0, -5.0])
y = np.array([1.0, -4.0, 8.0])
print_matrix("x", x)
print_matrix("y", y)

print_matrix("\|x\|", np.linalg.norm(x))
print_matrix("\|y\|", np.linalg.norm(y))
```

```
x
[[ 1. 10. -5.]
```

```
y
[[ 1. -4.  8.]
```

```
 $\|x\| = 11.224972160321824$ 
```

```
 $\|y\| = 9.0$ 
```




Different norms

- Euclidean norm or Euclidean distance measure

$$\|\mathbf{x} - \mathbf{y}\|_2$$

- L_p -norms or Minkowski norms, which is defined by:

$$\|\mathbf{x}\|_p = \left(\sum_i x_i^p \right)^{\frac{1}{p}}$$



$$L_2 = |A|$$


$$L_1 = |B| + |C|$$

$$L_\infty = \max(|B|, |C|) = |C|$$



Different norms


p	Notation	Common name	Effect	Uses	Geometric view
2	$ x $ or $ x _2$	Euclidean norm	Ordinary distance	Spatial distance measurement	Sphere just touching point
1	$ x _1$	Taxicab norm; Manhattan norm	Sum of absolute values	Distances in high dimensions, or on grids	Axis-aligned steps to get to point
0	$ x _0$	Zero pseudo-norm; non-zero sum	Count of non-zero values	Counting the number of "active elements"	Numbers of dimensions not touching axes
∞	$ x _{\text{inf}}$	Infinity norm; max norm	Maximum element	Capturing maximum "activation" or "excursion"	Smallest cube enclosing point
$-\infty$	$ x _{-\text{inf}}$	Min norm	Minimum element	Capturing minimum excursion	Distance of point to closest axis



Every dashed line has the **same** distance to the origin as measured in that norm. The points of equal distance in that norm appear as a connected line.

Unit vectors and normalisation

- A unit vector has norm 1 (the definition of a unit vector depends on the norm used)
 - Normalising for the **Euclidean norm** can be done by scaling the vector x by $\frac{1}{\|x\|_2}$
- If we think of vectors in the physics sense of having a **direction** and **length**, a unit vector is "**pure direction**".





Inner products of vectors

- An **inner product** $(\mathbb{R}^N \times \mathbb{R}^N) \rightarrow \mathbb{R}$ measures the **angle** between two real vectors.
 - It is related to the cosine distance:

$$\cos \theta = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}.$$

- For unit vectors, we can forget about the denominator, since $\|\mathbf{x}\| = 1, \|\mathbf{y}\| = 1$, so $\cos \theta = \mathbf{x} \cdot \mathbf{y}$.
- The computation of the **inner product**, for real-valued vectors in $\mathbb{R}^N$, is simply the sum of the elementwise products:

$$\mathbf{x} \cdot \mathbf{y} = \sum_i x_i y_i$$



Inner products of vectors

- Inner product in Numpy

```
x = np.array([1, 2, 3, 4])
y = np.array([4, 0, 1, 4])
print_matrix("x", x)
print_matrix("y", y)

print_matrix("x\cdot y", np.inner(x, y))
```

```
x
[[1 2 3 4]]
y
[[4 0 1 4]]
 $x \cdot y = 23$ 
```

- The inner product is only defined between vectors of the same dimension, and only in inner product spaces.
 - **ValueError**

- Don't mix it with the dot product

`numpy.dot(a, b, out=None)`

Dot product of two arrays. Specifically,

- If both a and b are 1-D arrays, it is inner product of vectors (without complex conjugation).
- If both a and b are 2-D arrays, it is matrix multiplication, but using `matmul` or `a @ b` is preferred.
- If either a or b is 0-D (scalar), it is equivalent to `multiply` and using `numpy.multiply(a, b)` or `a * b` is preferred.
- If a is an N-D array and b is a 1-D array, it is a sum product over the last axis of a and b .
- If a is an N-D array and b is an M-D array (where $M \geq 2$), it is a sum product over the last axis of a and the second-to-last axis of b :

For 2-D arrays it is the matrix product:

```
>>> a = [[1, 0], [0, 1]]
>>> b = [[4, 1], [2, 2]]
>>> np.dot(a, b)
array([[4, 1],
       [2, 2]])
```



Basic vector statistics

Some **statistics** that generalise the statistics of ordinary real numbers

- **mean vector** of a collection of N vectors

$$\text{mean}(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) = \frac{1}{N} \sum_i \mathbf{x}_i$$

- The mean vector is the **geometric centroid** of a set of vectors and can be thought of as capturing "centre of mass" of those vectors.

`numpy.mean` #

`numpy.mean(a, axis=None, dtype=None, out=None, keepdims=<no value>, *, where=<no value>)` [\[source\]](#)

`axis=0` for the purpose of mean vector

Examples

```
>>> a = np.array([[1, 2], [3, 4]])
>>> np.mean(a)
2.5
>>> np.mean(a, axis=0)
array([2., 3.])
>>> np.mean(a, axis=1)
array([1.5, 3.5])
```



Center a dataset


- We can **center** a dataset stored as an array of vectors to **zero mean** by just subtracting the mean vector from every row.

```
x_center = x - mu
print_matrix("x_center", x_center)
mu_c = np.mean(x_center, axis=0) #
print_matrix("\mu_c", mu_c)
```

```
x_center
[[-0.56  0.35 -1.3  -0.36]
 [-0.19  0.24  1.56  0.4 ]
 [ 0.75 -0.59 -0.27 -0.04]]
```

```
\mu c
```

```
[[ 0.  0.  0. -0.]]
```





High-dimensional vector spaces

- Data science often involves **high dimensional vector spaces**
 - High-dimensional can mean any $d > 3$;
 - a 20-dimensional feature set might be called medium-dimensional;
 - a 1000-dimensional might be called high-dimensional;
 - a 1M-dimensional dataset might be called extremely high-dimensional
- **Curse of dimensionality:** Many algorithms that work really well in low dimensions break down in higher dimensions.



Example: sailing weather station

- **Task:** to measure local atmospheric conditions during voyages
- **Input:** wind speed, temperature, humidity, sunshine hours, etc., over 10,000 measurements
- **Output:** *is it likely to be above 30C tomorrow?*




.Image by [Ronnierob](#) license [CC BY-SA](#)




Example: sailing weather station

Temperature (1D)



Now there are 20 bins in total for this dimension

Temperature & Humidity (3D)



Now there are 20 bins in each dimension, for 400 bins total.



Example: sailing weather station

- If we had 10 different measurements (air temperature, air humidity, latitude, longitude, wind speed, wind direction, precipitation, time of day, solar power, sea temperature)
 - we wanted to subdivide them into 20 bins each
 - How many bins in total?



Example: sailing weather station

- If we had 10 different measurements (air temperature, air humidity, latitude, longitude, wind speed, wind direction, precipitation, time of day, solar power, sea temperature)
 - We wanted to subdivide them into 20 bins each
 - How many bins in total?

We would need a histogram with 20^{10} bins --
over **10 trillion** bins.

- even using 8 bit unsigned integers this would be 10TB of memory
- But we only have 10,000 measurements



Example: sailing weather station

- If we had 10 different measurements (air temperature, air humidity, latitude, longitude, wind speed, wind direction, precipitation, time of day, solar power, sea temperature)
 - We wanted to subdivide them into 20 bins each
 - How many bins in total?

We would need a histogram with 20^{10} bins --
over **10 trillion** bins.

- even using 8 bit unsigned integers this would be 10TB of memory
- But we only have 10,000 measurements

Curse of dimensionality: as dimension increases generalisation gets harder *exponentially*



Matrices and linear operators

Matrices are 2D arrays of reals that define **linear maps**;

- Vectors represent “points in space”
- Matrices represent *operations* that do things to those points in space.

The *operations* represented by matrices are a particular class of *functions* on *vectors*

Operations with matrices

There are many things we can do with matrices:

- They can be **added** and **subtracted** $C = A + B$

$$(\mathbb{R}^{n \times m}, \mathbb{R}^{n \times m}) \rightarrow \mathbb{R}^{n \times m}$$

- They can be **scaled** with a scalar $C = sA$

$$(\mathbb{R}^{n \times m}, \mathbb{R}) \rightarrow \mathbb{R}^{n \times m}$$

- They can be **transposed** $B = A^T$; this exchanges rows and columns

$$\mathbb{R}^{n \times m} \rightarrow \mathbb{R}^{m \times n}$$

- They can be **applied** to vectors $\mathbf{y} = A\mathbf{x}$; this applies a matrix to a vector.

$$(\mathbb{R}^{n \times m}, \mathbb{R}^m) \rightarrow \mathbb{R}^n$$

- They can be **multiplied** together $C = AB$; this composes the effect of two matrices



Intro to matrix notation

- We write matrices as a capital letter, e.g. **A**:

$$A \in \mathbb{R}^{n \times m} = \begin{bmatrix} a_{1,1} & a_{1,2} & \dots & a_{1,m} \\ a_{2,1} & a_{2,2} & \dots & a_{2,m} \\ \dots & & & \\ a_{n,1} & a_{n,2} & \dots & a_{n,m} \end{bmatrix}, a_{i,j} \in \mathbb{R}$$

where each element of the matrix **A** is written as **A**_{*i,j*},
for the *i*th row and *j*th column.

```
a = np.array([[1, 2, 3],
              [4, 5, 6],
              [7, 8, 9]])

print_matrix("A", a)
# note that code indexes from 0

print_matrix("A_{1,3}", a[0,2])
```

A

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

$A_{1,3} = 3$



Matrices as maps

- We saw vectors as **points in space**, and matrices as **vector transform in space**.
- Matrices represent **linear maps** – these are functions applied to vectors, which output vectors. (Equivalent to applying some function $f(x)$ to the vectors)

$$A\mathbf{x} = f(\mathbf{x})$$

- A $n \times m$ matrix **A** represents a function $f(x)$ taking m dimensional vectors to n dimensional vectors ($\mathbb{R}^m \rightarrow \mathbb{R}^n$)
- Matrices capture a special set of functions that preserve important properties of the vectors they act on.



Linear maps

Linearity

- the transform of the sum of two vectors is the **same** as the sum of the transform of two vectors
- the transform of a scalar multiple of a vector is the **same** as the scalar multiple of the transform of a vector

$$\begin{aligned} f(\mathbf{x} + \mathbf{y}) &= f(\mathbf{x}) + f(\mathbf{y}) &= A(\mathbf{x} + \mathbf{y}) &= A\mathbf{x} + A\mathbf{y}, \\ f(c\mathbf{x}) &= cf(\mathbf{x}) &= A(c\mathbf{x}) &= cA\mathbf{x}, \end{aligned}$$

Anything which is linear is easy. Anything which isn't linear is hard.



Transforms and projections

- A **linear map** is any function $f: R^m \rightarrow R^n$ which satisfies the **linearity** requirements.
- If the map represented by the matrix is $n \times n$ then it maps from a vector space onto the **same** vector space (e.g. from $\mathbb{R}^n \rightarrow \mathbb{R}^n$), and it is called a **linear transform**.
- If the map has the property $Ax = AAx$ or equivalently $f(x) = f(f(x))$ then the operation is called a **linear projection**;
 - for example, projecting 3D points onto a plane; applying this transform to a set of vectors twice is the same as applying it once.


Matrices represent linear maps or linear functions.



Examples - effect of linear transforms

- linear transforms (linear maps $A \in \mathbb{R}^{2 \times 2}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$), on the 2D plane.
- We form the product Ax , which "applies" the matrix to the vector x .


Identity
[[1 0]
[0 1]]



Examples - effect of linear transforms

- linear transforms (linear maps $A \in \mathbb{R}^{2 \times 2}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$), on the 2D plane.
- We form the product Ax , which "applies" the matrix to the vector x .

Uniform scale
[[0.5 0.]
[0. 0.5]]






Examples - effect of linear transforms

- linear transforms (linear maps $A \in \mathbb{R}^{2 \times 2}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$), on the 2D plane.
- We form the product Ax , which "applies" the matrix to the vector x .

Non-uniform scale


```
[[0.5  0. ]  
[0.   1. ]]
```



Examples - effect of linear transforms

- linear transforms (linear maps $A \in \mathbb{R}^{2 \times 2}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$), on the 2D plane.
- We form the product Ax , which "applies" the matrix to the vector x .

Rotate 90
[[0 1]
[-1 0]]




Examples - effect of linear transforms

- linear transforms (linear maps $A \in \mathbb{R}^{2 \times 2}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$), on the 2D plane.
- We form the product Ax , which "applies" the matrix to the vector x .

Rotate 30

$$\begin{bmatrix} 0.87 & 0.5 \\ -0.5 & 0.87 \end{bmatrix}$$






Examples - effect of linear transforms

- linear transforms (linear maps $A \in \mathbb{R}^{2 \times 2}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$), on the 2D plane.
- We form the product Ax , which "applies" the matrix to the vector x .

Rotate 45, Scale 0.5
[[0.35 0.35]
[-0.35 0.35]]





Matrix operations

The addition of matrices of *equal size* is simple elementwise addition.

$$A + B = \begin{bmatrix} a_{1,1} + b_{1,1} & a_{1,2} + b_{1,2} & \dots & a_{1,m} + b_{1,m} \\ a_{2,1} + b_{2,1} & a_{2,2} + b_{2,2} & \dots & a_{2,m} + b_{2,m} \\ \dots & \dots & \dots & \dots \\ a_{n,1} + b_{n,1} & a_{n,2} + b_{n,2} & \dots & a_{n,m} + b_{n,m} \end{bmatrix}$$



Matrix operations

The scalar multiplication cA is to multiply each element by c .

$$cA = \begin{bmatrix} ca_{1,1} & ca_{1,2} & \dots & ca_{1,m} \\ ca_{2,1} & ca_{2,2} & \dots & ca_{2,m} \\ \dots & & & \\ ca_{n,1} & ca_{n,2} & \dots & ca_{n,m} \end{bmatrix}$$



Application to vectors

- We can apply a matrix to a vector, equivalent to applying the function $f(x)$ to x .

$$A\mathbf{x} = f(\mathbf{x})$$

- If A is $\mathbb{R}^{n \times m}$, and $\mathbf{x}$ is $\mathbb{R}^{m \times 1}$, then this will map from an m dimensional vector space to an n dimensional vector space: $(\mathbb{R}^{n \times m}, \mathbb{R}^m) \rightarrow \mathbb{R}^n$.
- All application of a matrix to a vector does is form a
 - **weighted sum of the elements of the vector.**
- This is a linear combination (equivalent to a "weighted sum") of the components.



Application to vectors

In particular, we take each element of $\mathbf{x}$, $x_1, x_2, \dots, x_m$, multiply it with the corresponding column of A , and sum these columns together.

- Set $\mathbf{y} = [0, 0, 0, \dots] = \mathbf{0}^n$ (the n -dimensional zero vector)
- For each column $1 \leq i \leq m$ in A
- $\mathbf{y} = \mathbf{y} + x_i A_i$. Note that $x_i A_i$ is scalar times vector, and has n elements. A_i here means the i th column of A .



University
of Glasgow

Let's have a try

$$f(x)=Ax$$

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \text{ and } x = \begin{bmatrix} 1 \\ 2 \end{bmatrix} Ax = ?$$

Let's have a try

$$f(x)=Ax$$

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \text{ and } x = \begin{bmatrix} 1 \\ 2 \end{bmatrix} Ax = ?$$

The traditional matrix multiplication approach (row-by-column)

The matrix-vector product is:

$$Ax = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 1(1) + 2(2) \\ 3(1) + 4(2) \\ 5(1) + 6(2) \\ 7(1) + 8(2) \end{bmatrix} = \begin{bmatrix} 5 \\ 11 \\ 17 \\ 23 \end{bmatrix} \text{ So, } Ax = \begin{bmatrix} 5 \\ 11 \\ 17 \\ 23 \end{bmatrix}.$$

Let's have a try

$$f(x)=Ax$$

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \text{ and } x = \begin{bmatrix} 1 \\ 2 \end{bmatrix} Ax = ?$$

An alternative way :

1. Multiply the first element of (x) with the first column of (A):
 $1 \times [1, 3, 5, 7] = [1, 3, 5, 7]$
2. Multiply the second element of (x) with the second column of (A):
 $2 \times [2, 4, 6, 8] = [4, 8, 12, 16]$
3. Sum the results from steps 1 and 2 element-wise:
 $[1, 3, 5, 7] + [4, 8, 12, 16] = [5, 11, 17, 23]$

So, the result is the same as before: ($[5, 11, 17, 23]$).



Application to vectors

- We can use `@` to form products of vectors and matrices in Numpy

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} \text{ and } x = \begin{bmatrix} 1 \\ 2 \end{bmatrix} Ax = ?$$

```
A = np.array([[1, 2], [3, 4], [5, 6], [7, 8]])
x = np.array([1, 2])

print_matrix("A", A)
print_matrix("\bf x", x)
print_matrix("A\\bf x", A @ x)
```

```
A
[[1 2]
 [3 4]
 [5 6]
 [7 8]]
\bf x
[[1 2]]
A\\bf x
[[ 5 11 17 23]]
```




Matrix multiplication

- **Matrix multiplication** defines the product $\mathbf{C}=\mathbf{AB}$, where $\mathbf{A}, \mathbf{B}, \mathbf{C}$ are all matrices.
 - Matrix multiplication is defined such that if $\mathbf{A}$ represents linear transform $f(\mathbf{x})$ and $\mathbf{B}$ represents linear transform $g(\mathbf{x})$, then $\mathbf{BAx}=g(f(\mathbf{x}))$
 - **Multiplying two matrices is equivalent to composing the linear functions they represent, and it results in a matrix which has that affect.**
-
- *Note that the composition of linear maps is read right to left. To apply the transformation $\mathbf{A}$, then $\mathbf{B}$, we form the product $\mathbf{BA}$, and so on.*



Multiplication algorithm

- If $C=AB$ then

$$C_{ij} = \sum_k a_{ik} b_{kj}$$

- **Multiplication** is *only* defined for two matrices **A**,**B** if:
 - **A** is $p \times q$ and
 - **B** is $q \times r$.
- This gives rise to many important uses of matrices, for example, the product of a scaling matrix and a rotation matrix is a scale-and-rotate matrix.



Multiplication in Python

- Matrix multiplication is of course built into NumPy
- Matrix multiplication is applied by `np.dot(a, b)`
 - or – better – by the syntax `a @ b`

```
# verify that this is the same as the built-in  
c_numpy = np.dot(a, b)  
print_matrix("C_{\\text numpy}", c_numpy)  
print(np.allclose(c, c_numpy))
```

```
c_at = a @ b  
print_matrix("C_{\\text a @ b}", a @ b)  
print(np.allclose(c_at, c))
```

```
C_{\\text numpy}
```

```
[[ -4   4  -4]]
```

```
True
```

```
C_{\\text a @ b}
```

```
[[ -4   4  -4]]
```

```
True
```



Commutativity and Transpose

- The **order** of multiplication is important.
- Matrix multiplication does **not** commute

$$AB \neq BA$$

- Transpose order switching

$$(AB)^T = B^T A^T$$

- It is also true that

$$(A + B)^T = A^T + B^T$$





Time complexity of multiplication

- Matrix multiplication has, in the general case, of time complexity $O(pqr)$, or for multiplying two square matrices $O(n^3)$.
- However, there are many special forms of matrices for which this complexity can be reduced, such as *diagonal*, *triangular*, *sparse* and *banded* matrices. (later)
- There are some accelerated algorithms for general multiplication. The time complexity of all of them is $>O(N^2)$ but $<O(N^3)$.



Transposition

- The **transpose** of a matrix A is written A^T and has the same elements, but with the rows and columns exchanged.
- Two ways of using numpy
 - $A.T$
 - `np.transpose(A)`

```
### Transpose
```

```
A = np.array([[2, -5], [1, 0], [3, 3]])  
print_matrix("A", A)  
print_matrix("A^T", A.T)
```

A

```
[[ 2 -5]  
 [ 1  0]  
 [ 3  3]]
```

A^T

```
[[ 2  1  3]  
 [-5  0  3]]
```

```
>>> a = np.array([[1, 2], [3, 4]])  
>>> a  
array([[1, 2],  
       [3, 4]])  
>>> np.transpose(a)  
array([[1, 3],  
       [2, 4]])
```



Special matrix multiplication

- **outer product:** the product of a $M \times 1$ with a $1 \times N$ vector, which is an $M \times N$ matrix

$$\mathbf{x} \otimes \mathbf{y} = \mathbf{x}^T \mathbf{y}$$

- **inner product:** the product of a $1 \times N$ with an $N \times 1$ vector is a 1×1 matrix, which is a scalar

$$\mathbf{x} \bullet \mathbf{y} = \mathbf{x} \mathbf{y}^T$$




Matrix multiplication as composed maps

- We saw vectors as **points in space**, and matrices as **vector transform in space**.
- **Multiplication is composition:** If A represents $f(x)$ and B represents $g(x)$, then the product BA represents $g(f(x))$.
- $BAx = B(Ax)$ means do A to x , then do B to the result




Composed maps

- nonuniform scaling



- rotation





Composed maps

- Nonuniform scaling matrix: `scale_x`
- Rotation matrix: `rot30`

$$\text{Rotate_then_scale} = \text{scale_x} \times \text{rot30}$$


$$= \begin{bmatrix} 1 & 0 \\ 0 & 0.5 \end{bmatrix} \times \begin{bmatrix} 0.87 & 0.5 \\ -0.5 & 0.87 \end{bmatrix}$$

$$= \begin{bmatrix} 0.87 & 0.5 \\ -0.25 & 0.435 \end{bmatrix}$$

Rotate then scale

```
[[ 0.87  0.5 ]  
 [-0.25  0.43]]
```

Rotate then scale



Composed maps

- Nonuniform scaling matrix: `scale_x`
- Rotation matrix: `rot30`

$$\text{Scale_then_rotate} = \text{rot30} \times \text{scale_x}$$


$$= \begin{bmatrix} 0.87 & 0.5 \\ -0.5 & 0.87 \end{bmatrix} \times \begin{bmatrix} 1 & 0 \\ 0 & 0.5 \end{bmatrix}$$

$$= \begin{bmatrix} 0.87 & 0.25 \\ -0.5 & 0.435 \end{bmatrix}$$

Scale then rotate

```
[[ 0.87  0.25]
 [-0.5   0.43]]
```

Scalethenrotate





Concatenation of transforms

- Many software operations take advantage of the definition of matrix multiplication as the composition of linear maps.
- In a graphics processing pipeline, for example, all of the operations to position, scale and orient visible objects are represented as matrix transforms.
- Multiple operations can be combined into *one single matrix operation*.



An example matrix for measuring spread: covariance matrices

- **variance:** measures the spread of a dataset

$$\sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2$$

- ***Covariance between two datasets / populations:*** measures how they co-vary

$$Cov(X, Y) = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$$

- $Cov > 0$: Variable variances coincides – as one increases so does the other
- $Cov < 0$: Opposite
- $Cov = 0$: No linear relationship between X and Y
- Cov is not standardised – its magnitude is not bounded, it depends on Xs and Ys.



An example matrix for measuring spread: Covariance Matrices

- If you have multiple variables $X_1 \dots X_n$
- The covariance matrix is

$$\Sigma = \text{Cov}(\mathbf{X}) = \begin{bmatrix} \text{Cov}(X_1, X_1) & \text{Cov}(X_1, X_2) & \cdots & \text{Cov}(X_1, X_n) \\ \text{Cov}(X_2, X_1) & \text{Cov}(X_2, X_2) & \cdots & \text{Cov}(X_2, X_n) \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}(X_n, X_1) & \text{Cov}(X_n, X_2) & \cdots & \text{Cov}(X_n, X_n) \end{bmatrix}$$

- The main diagonal has the variances of individual populations or variables
- Symmetric
- Eigenvalues are non-negative (important for some applications, see later)

An example matrix for measuring spread: covariance matrices

- Generate data examples

```
x = np.random.normal(0,1,(500, 5))

mu = np.mean(x, axis=0)
sigma_cross = ((x - mu).T @ (x - mu)) / (x.shape[0]-1)
np.set_printoptions(suppress=True, precision=2)
print_matrix("\Sigma_{\text{cross}}", sigma_cross)
```

```
\Sigma_{\text{cross}}
[[ 1.04  0.04  0.02  0.04 -0.02]
 [ 0.04  1.03 -0.03  0.07 -0.03]
 [ 0.02 -0.03  1.04 -0.08  0.01]
 [ 0.04  0.07 -0.08  1.09 -0.02]
 [-0.02 -0.03  0.01 -0.02  0.98]]
```

- 500x5 matrix

- Compute covariance using `np.cov()`

```
# verify it is close to the function provided by NumPy
sigma_np = np.cov(x, rowvar=False)
print_matrix("\Sigma_{\text{numpy}}",
            sigma_np)
```


```
\Sigma_{\text{numpy}}
[[ 1.04  0.04  0.02  0.04 -0.02]
 [ 0.04  1.03 -0.03  0.07 -0.03]
 [ 0.02 -0.03  1.04 -0.08  0.01]
 [ 0.04  0.07 -0.08  1.09 -0.02]
 [-0.02 -0.03  0.01 -0.02  0.98]]
```

- 5x5 matrix



Covariance ellipses

- For a 2-D dataset
 - **mean vector: 1×2**
 - **Covariance matrix: 2×2**
- The mean vector captures the idea of "centre"
- The covariance matrix captures the "spread" of a collection of points in a vector space.





Special matrix forms

- **Diagonal matrix**

- `np.diag(x)`

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 3 \end{bmatrix}$$

- **Identity matrix I**

- `np.eye(n)`

$$\begin{bmatrix} 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \dots & & & & \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix}$$

- **Zero matrix**

- `np.zeros()`

$$\begin{bmatrix} [0. & 0. & 0. & 0.] \\ [0. & 0. & 0. & 0.] \\ [0. & 0. & 0. & 0.] \\ [0. & 0. & 0. & 0.] \end{bmatrix}$$

- **Upper triangular**

$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 0 & 5 & 6 & 7 \\ 0 & 0 & 8 & 9 \\ 0 & 0 & 0 & 10 \end{bmatrix}$$

- **Lower triangular**

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 2 & 3 & 0 & 0 \\ 4 & 5 & 6 & 0 \\ 7 & 8 & 9 & 10 \end{bmatrix}$$



Matrix Inversion

Four basic algebraic operations on matrices:

- scalar multiplication $c\mathbf{A}$;
- matrix addition $\mathbf{A}+\mathbf{B}$;
- matrix multiplication $\mathbf{BA}$
- matrix transposition $\mathbf{A}^T$

A further important operation: **inversion** $\mathbf{A}^{-1}$, defined such that $\mathbf{A}^{-1}\mathbf{A}=\mathbf{I}$:

- $\mathbf{A}^{-1}(\mathbf{Ax})=\mathbf{x}$,
- $(\mathbf{A}^{-1})^{-1}=\mathbf{A}$
- $(\mathbf{AB})^{-1}=\mathbf{B}^{-1}\mathbf{A}^{-1}$

*Inversion is only defined for certain kinds of matrices: **Square Matrices, Non-singular Matrices***

- Inversion of a matrix creates a new linear operator which reverses the original operation. Let's see it in action:

Inversion as "undo"

```
# Verify that left-multiplying by the inverse "undoes" the transformation applied by A to a random vector
# Create a random vector
x = np.random.normal(0, 1, (3, 1))
print_matrix("x", x)
# Transform it
print_matrix("Ax", A @ x)
# Left-multiply by the inverse to recover the original vector
print_matrix("A^{-1}(Ax)", np.linalg.inv(A) @ (A @ x))
```

```
x
[[-0.18]
 [ 0.95]
 [-0.92]]
```


```
Ax
[[-0.28]
 [ 0.73]
 [-1.11]]
```

```
A^{-1}(Ax)
[[-0.18]
 [ 0.95]
 [-0.92]]
```




Inversion as "undo"

- 30 degree rotation matrix



- 30 degree rotation matrix, left-multiplied by its inverse





Beyond this course

- 3blue1brown Linear Algebra series (strongly recommended)
- Introduction to applied linear algebra by *S. Boyd and L. Vandenberghe*
- **Linear Algebra Done Right** by *Sheldon Axler* (excellent introduction to the "pure math" side of linear algebra) ISBN-13: 978-0387982588
- **Coding the Matrix: Linear Algebra through Applications to Computer Science** by *Philip N Klein* (top quality textbook on how linear algebra is implemented, all in Python) ISBN-13: 978-0615880990
- **Linear Algebra and Learning from Data** *Gilbert Strang*, ISBN-13: 978-069219638-0, explains many detailed aspects of linear algebra and how they relate to data science.
- The Matrix Cookbook by *Kaare Brandt Petersen* and *Michael Syskind Pedersen*. If you need to do a tricky calculation with matrices, this book will probably tell you how to do it.



University
of Glasgow

Thank you

Contact:
iraklis.klampanos@glasgow.ac.uk